

CSI 2300: Intro to Data Science

In-Class Exercise 22: Modeling – Nonlinear Regression

For this lecture, we are going to work with a Shiny app built for you.

1. Load the file `modeling_22_shiny_nonlinear.R` into RStudio, and get it running.

Getting the app running properly on all datasets requires the `mowaterR` and `fields` libraries, so make sure to install those first. If you cannot get the app running on your computer, that's okay, so long as someone in your group can, and you participate.

In this app, you can select among several datasets. For each dataset, there is a fixed y variable (the response, or dependent variable). You can select among several x variables (the predictor, or independent variable). You can also select to use log-scaling on either (or both) axes. Try playing around with these different options (dataset, predictor, log-scaling) to get a feel for the different datasets.

2. Your group should have an assigned dataset and independent variable. For that data, do the following:
 - Look at the scatter plot (observations, not the fitted curves). Does there appear to be a relationship between x and y ?
 - If there is a relationship, is it easier to observe that relationship with the original data, or with log-scale (x , y , or both)?
3. The parameters allow a fairly wide amount of control over the “flexibility” of the nonlinear models. The `loess` and `lowess` models have the `span` and `f` parameters, while the `polynomial` model has the highest polynomial degree (e.g. the largest k in x^k), and `LASSO` has the penalty term λ (lambda). Here, `LASSO` is using a penalized 20th-order polynomial, while the `polynomial` model limits the largest polynomial degree.
 - For each of the four curves, adjusting its main parameter increases or decreases its flexibility. Which direction is which for each of these curves (e.g. “increasing/decreasing the span for loess increases/decreases its flexibility...”)?
 - What happens to how well the curves fit when you go too far in either direction (too much flexibility, not enough flexibility)? (There may be extreme cases where the software fails to work properly, that's just your professor's poor code, don't report on that.)
4. We have two types of models shown here: non-parametric (`loess` and `lowess`), and parametric (`polynomial` and `LASSO`). What do you observe that is different about how the two types of models behave (locally, globally) and are able (or not able) to fit the data?
5. Now, let's investigate the quality of the model fits.

- Each of the four curve types has a parameter you can adjust. Adjust each one's parameter to obtain a fit to the data that **visually captures the overall pattern**. Report the parameter that works best for each curve and the corresponding R^2 value.
 - Now adjust each curve's parameter to obtain a fit to the data that has **the highest R^2** . Report the parameter and the corresponding R^2 value. Is this a good fit to the overall pattern of the data?
 - The results of this will be reported back to the whole class, so be prepared to share your screen with your two sets of best fit curves (best overall pattern and highest R^2) and explain them.
 - In your report, show a screenshot of your best model fits.
6. We know of a couple of different ways of measuring model fitness. The R^2 measure is the most widely known. But we know that by itself, R^2 isn't great – we can get a high R^2 but a poor model fit by overfitting, or “memorizing”, the data used for fitting the model.

Suggest a method by which we could get a *more reliable* R^2 estimate than the one that is based on the observations used to fit the model. Think about using data that was not used for fitting the model.